

CIND820

November 9, 2021

```
[1]: import pandas as pd
from matplotlib import pyplot as plt
import numpy as np
import seaborn as sns
import numpy as np
df=pd.read_csv("Heart_Disease.csv")#Load the data into the system
df.info() #determine the attribute and data type
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   303 non-null    int64
1   Sex                                   303 non-null    int64
2   Chest_Pain                           303 non-null    int64
3   Resting_Blood_Pressure               303 non-null    int64
4   Colestrol                            303 non-null    int64
5   Fasting_Blood_Sugar                  303 non-null    int64
6   Rest_ECG                             303 non-null    int64
7   MAX_Heart_Rate                       303 non-null    int64
8   Exercised_Induced_Angina             303 non-null    int64
9   ST_Depression                        303 non-null    float64
10  Slope                                303 non-null    int64
11  Major_Vessels                         303 non-null    object
12  Thalessemia                          303 non-null    object
13  Target                                303 non-null    int64
dtypes: float64(1), int64(11), object(2)
memory usage: 33.3+ KB
```

```
[2]: df.head(10) #showing the first 10 rows
```

```
[2]:   Age  Sex  Chest_Pain  Resting_Blood_Pressure  Colestrol  \
0   63   1         1             145           233
1   67   1         4             160           286
2   67   1         4             120           229
3   37   1         3             130           250
4   41   0         2             130           204
```

5	56	1	2	120	236
6	62	0	4	140	268
7	57	0	4	120	354
8	63	1	4	130	254
9	53	1	4	140	203

	Fasting_Blood_Sugar	Rest_ECG	MAX_Heart_Rate	Exercised_Induced_Angina	\
0	1	2	150		0
1	0	2	108		1
2	0	2	129		1
3	0	0	187		0
4	0	2	172		0
5	0	0	178		0
6	0	2	160		0
7	0	0	163		1
8	0	2	147		0
9	1	2	155		1

	ST_Depression	Slope	Major_Vessels	Thalessemia	Target
0	2.3	3	0	6	0
1	1.5	2	3	3	2
2	2.6	2	2	7	1
3	3.5	3	0	3	0
4	1.4	1	0	3	0
5	0.8	1	0	3	0
6	3.6	3	2	3	3
7	0.6	1	0	3	0
8	1.4	2	1	7	2
9	3.1	3	0	7	1

```
[3]: #checking for missing value, replacing '?' to NAN
df.replace("?", np.nan, inplace = True)
df.isnull().sum()
```

```
[3]: Age          0
Sex             0
Chest_Pain      0
Resting_Blood_Pressure  0
Colestrol       0
Fasting_Blood_Sugar  0
Rest_ECG        0
MAX_Heart_Rate  0
Exercised_Induced_Angina  0
ST_Depression   0
Slope           0
Major_Vessels   4
Thalessemia     2
```

```
Target          0
dtype: int64
```

```
[4]: #drop the rows that contain missing values
df.dropna(inplace= True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 297 entries, 0 to 301
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    297 non-null    int64
1   Sex                                    297 non-null    int64
2   Chest_Pain                            297 non-null    int64
3   Resting_Blood_Pressure                 297 non-null    int64
4   Colestrol                              297 non-null    int64
5   Fasting_Blood_Sugar                    297 non-null    int64
6   Rest_ECG                               297 non-null    int64
7   MAX_Heart_Rate                         297 non-null    int64
8   Exercised_Induced_Angina               297 non-null    int64
9   ST_Depression                          297 non-null    float64
10  Slope                                  297 non-null    int64
11  Major_Vessels                          297 non-null    object
12  Thalessemia                            297 non-null    object
13  Target                                  297 non-null    int64
dtypes: float64(1), int64(11), object(2)
memory usage: 34.8+ KB
```

```
[5]: #change the target value to 0 and 1
df['Target'] = df.Target.apply(lambda q: '0' if q == 0 else '1')
print(df)
df['Target'].value_counts()
```

	Age	Sex	Chest_Pain	Resting_Blood_Pressure	Colestrol	\
0	63	1	1	145	233	
1	67	1	4	160	286	
2	67	1	4	120	229	
3	37	1	3	130	250	
4	41	0	2	130	204	
..	
297	57	0	4	140	241	
298	45	1	1	110	264	
299	68	1	4	144	193	
300	57	1	4	130	131	
301	57	0	2	130	236	

	Fasting_Blood_Sugar	Rest_ECG	MAX_Heart_Rate	Exercised_Induced_Angina	\
--	---------------------	----------	----------------	--------------------------	---

0	1	2	150	0
1	0	2	108	1
2	0	2	129	1
3	0	0	187	0
4	0	2	172	0
..
297	0	0	123	1
298	0	0	132	0
299	1	0	141	0
300	0	0	115	1
301	0	2	174	0

	ST_Depression	Slope	Major_Vessels	Thalessemia	Target
0	2.3	3	0	6	0
1	1.5	2	3	3	1
2	2.6	2	2	7	1
3	3.5	3	0	3	0
4	1.4	1	0	3	0
..
297	0.2	2	0	7	1
298	1.2	2	0	7	1
299	3.4	2	2	7	1
300	1.2	2	1	7	1
301	0.0	2	1	3	1

[297 rows x 14 columns]

```
[5]: 0    160
      1    137
      Name: Target, dtype: int64
```

```
[6]: print(df.dtypes)
```


```
Age                int64
Sex                int64
Chest_Pain         int64
Resting_Blood_Pressure  int64
Colestrol          int64
Fasting_Blood_Sugar  int64
Rest_ECG           int64
MAX_Heart_Rate     int64
Exercised_Induced_Angina  int64
ST_Depression      float64
Slope              int64
Major_Vessels      object
Thalessemia        object
Target             object
dtype: object
```

```
[7]: #changing the data types to interger
df = df.apply(pd.to_numeric)
```

```
[8]: print(df.dtypes)
```


```
Age                int64
Sex                int64
Chest_Pain         int64
Resting_Blood_Pressure  int64
Colestrol          int64
Fasting_Blood_Sugar  int64
Rest_ECG           int64
MAX_Heart_Rate     int64
Exercised_Induced_Angina  int64
ST_Depression      float64
Slope              int64
Major_Vessels      int64
Thalessemia        int64
Target             int64
dtype: object
```

```
[9]: #using heatmap to see the correlation between all the variable and pick the
      ↳ variable that has the closer relationship with the targer value.
      #using this method to pick the significant features to test our classificaiton
      ↳ model
      #base on the result 4 variable have been kicked out which are
      ↳ "Resting_Blood_Pressure", "Colestrol", "MAX_Heart_Rate", and
      ↳ "Fasting_Blood_Sugar".
corrmat = df.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20))
g=sns.heatmap(df[top_corr_features].corr(),annot=True)
```




```
[10]: #plot all the distribution of 9 variables against target variable.
sns.countplot(x='Age', data=df)
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e434b4690>
```




```
[11]: sns.countplot(x='Sex', data=df)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43e464d0>
```



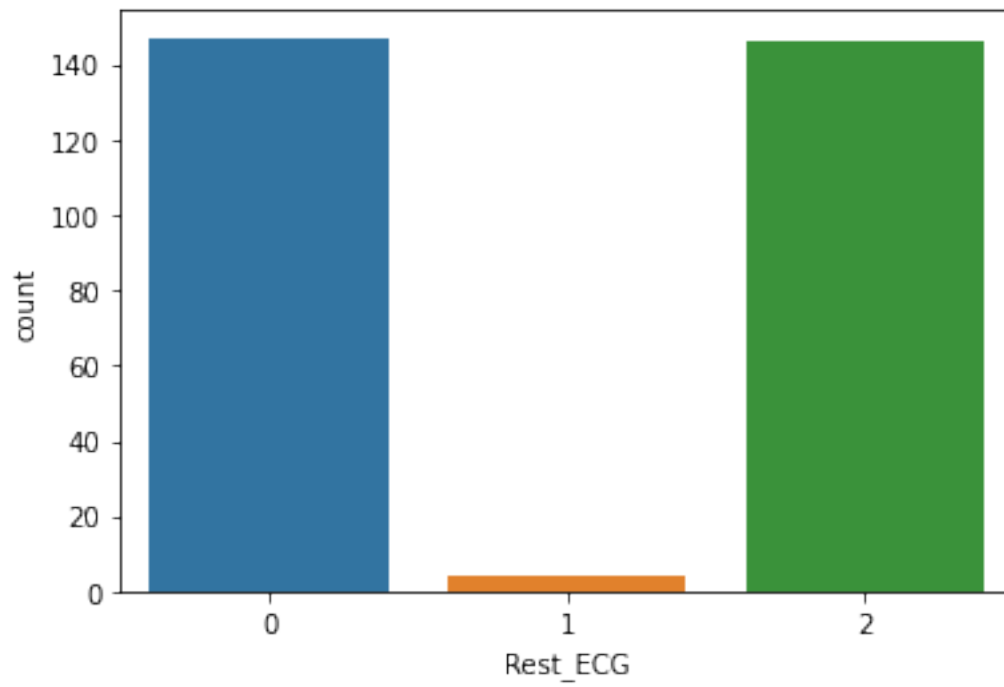
```
[12]: sns.countplot(x='Chest_Pain', data=df)
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43e15710>
```



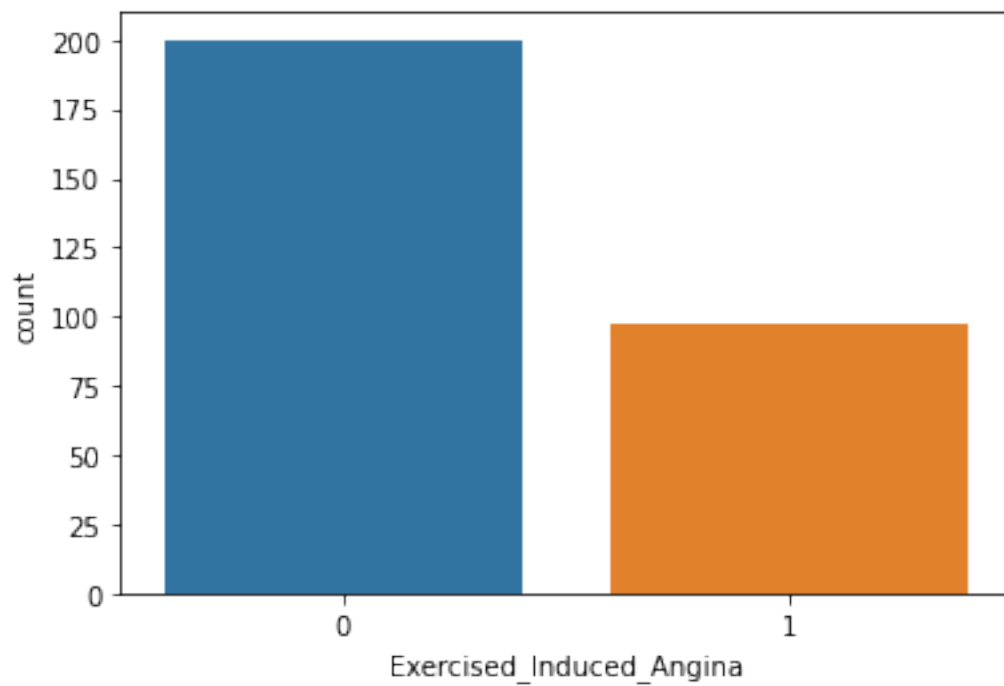
```
[13]: sns.countplot(x='Rest_ECG', data=df)
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43d8cc10>
```


```
[14]: sns.countplot(x='Exercised_Induced_Angina', data=df)
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43ea1e10>
```




```
[15]: sns.countplot(x='ST_Depression', data=df)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43e1cf50>
```



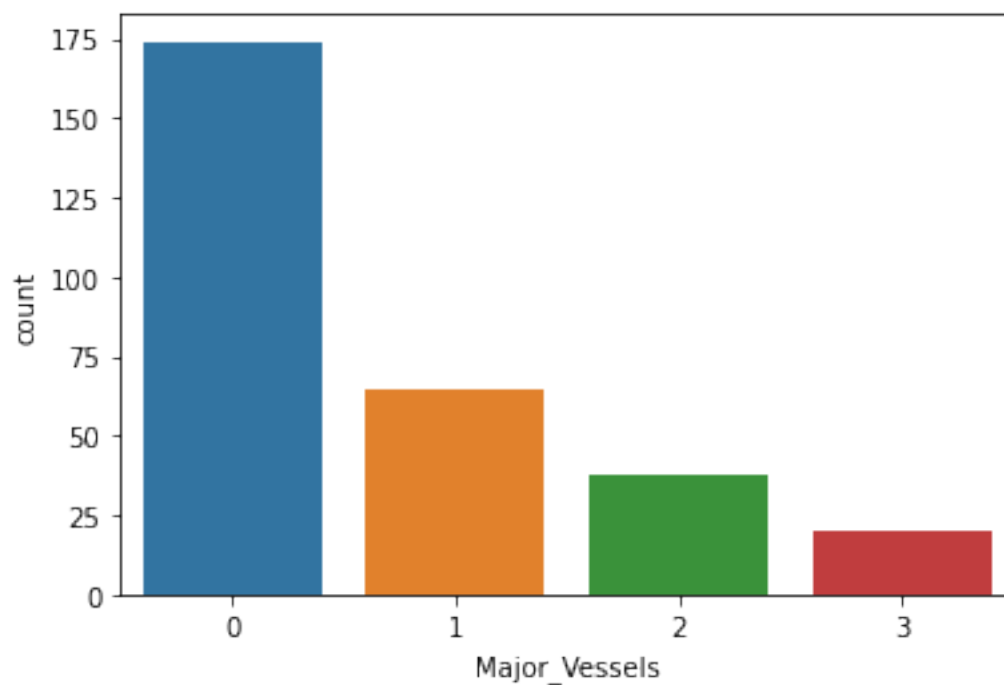
```
[16]: sns.countplot(x='Slope', data=df)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43b65950>
```



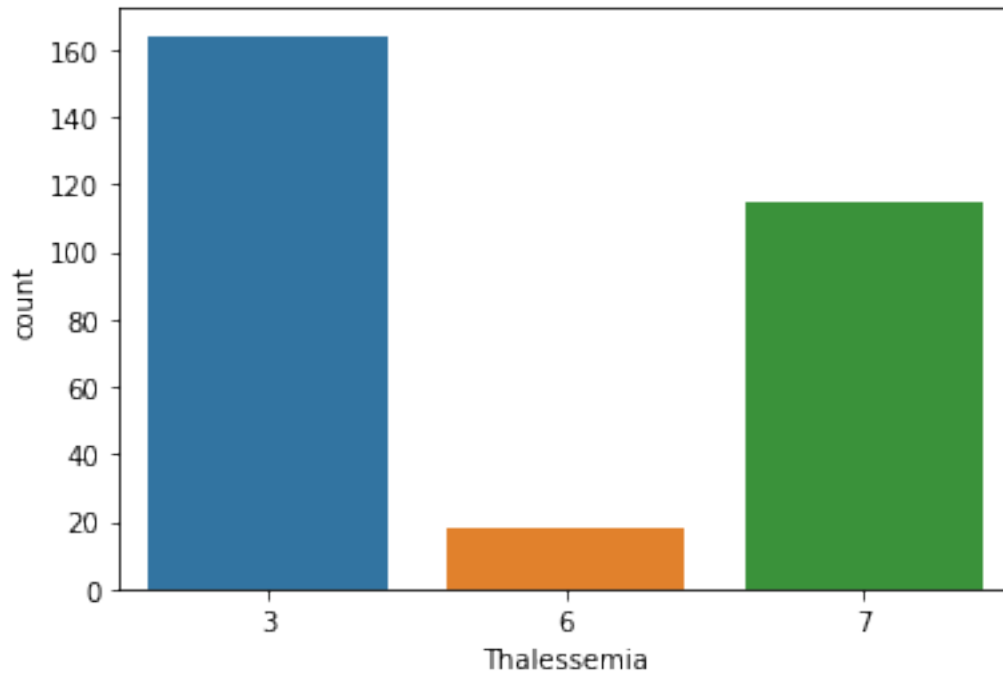
```
[17]: sns.countplot(x='Major_Vessels', data=df)
```

```
[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43b41050>
```



```
[18]: sns.countplot(x='Thalessemia', data=df)
```

```
[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7e43aada90>
```



```
[21]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

feature_cols = ['Age', 'Sex', 'Rest_ECG',
    ↳ 'Chest_Pain', 'Exercised_Induced_Angina', 'ST_Depression', 'Slope', 'Major_Vessels', 'Thalessemi
X = df[feature_cols] # Features
y = df.Target
```

```
[22]: #setting up training set and testing set with 0.25 size
#this result in 74 vs 223
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
    ↳ 25,random_state=0)
```

```
[23]: #using logic regression algorithm to find the accuracy
logreg = LogisticRegression()
```

```
logreg.fit(X_train,y_train)
y_pred=logreg.predict(X_test)
```

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
[24]: from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix
```

```
[24]: array([[37,  2],
          [10, 26]])
```

```
[25]: print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))
```

Accuracy: 0.84
Precision: 0.9285714285714286
Recall: 0.7222222222222222

```
[26]: #using KNN algorithm model to find the accuracy

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=25)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))
```

Accuracy: 0.8266666666666667
Precision: 0.896551724137931
Recall: 0.7222222222222222

```
[27]: #using decision tree algorithm model to find the accuracy

from sklearn.tree import DecisionTreeClassifier # Import Decision Tree_
      ↪Classifier
```

```
clf = DecisionTreeClassifier()
clf = clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))
```

Accuracy: 0.76
Precision: 0.75
Recall: 0.75

[28]: *#using naive bayes algorithm model to find the accuracy*

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred = gnb.predict(X_test)
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))
```

Accuracy: 0.8533333333333334
Precision: 0.9310344827586207
Recall: 0.75